

Introduction



This dataset is designed for **regression analysis** and is mainly used for **machine learning practice**. It contains **numerical features only**, which makes it suitable for applying different regression models without heavy preprocessing.

The dataset consists of **10,000 rows** and **7 columns**, where **6 columns are input features**: **1 column is the target variable**. There are **no missing values**, which helps in faster and cleaner model training.

Import libarays

In [6]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from IPython.display import HTML
import warnings
warnings.filterwarnings("ignore")

plt.rcParams["figure.figsize"] = (10, 6)
print('All libraries imported successfully')
```

All libraries imported successfully



Loading the Dataset



```
In [7]: data = pd.read_csv("practice_dataset.csv")
data.head()
```

```
Out[7]:
```

	a	b	c	d	e	f	target
0	3.745401	7.472816	2.299983	6.743301	-4.021759	2.541710	-23.804962
1	9.507143	6.658242	-3.154880	5.133632	-8.103644	1.483551	-158.554897
2	7.319939	3.523078	-1.533603	9.680487	-7.472816	0.586397	-69.630480
3	5.986585	12.145333	1.632806	2.970806	-6.386577	2.209925	-77.556230
4	1.560186	9.532483	-0.179107	6.290708	-5.926933	1.256034	-93.973154

Exploratory Data Analysis(EDA)

```
In [8]: # -----
print("*"*70)
print("\033[1m" + "Columns Names".center(70) + "\033[0m")
print("*"*70)
practice = pd.DataFrame(data.columns, columns=["Columns"])
display(practice)
# -----
print("*"*70)
#print("")
print("\033[1m" + "Information Dataset ".center(70) + "\033[0m")
print("*"*70)
print(" ")
# -----
data.info()
print("*"*70)
print("\033[1m" + "Data Types from Dataset".center(70) + "\033[0m")
print("*"*70)
display(data.dtypes)
# -----
print("*"*70)
print("\033[1m" + "Description".center(70) + "\033[0m")
print("*"*70)
display(data.describe())
# -----
print("*"*70)
print("\033[1m" + "Missing Values".center(70) + "\033[0m")
print("*"*70)
```

```

display(data.isnull().sum())
# -----
print("*"*70)
print("\033[1m" + "Duplicated Values".center(70) + "\033[0m")
print("*"*70)
display(data.duplicated())
# -----
print("*"*70)
print("\033[1m" + "Shape of the Dataset ".center(70) + "\033[0m")
print("*"*70)
var_1 = pd.DataFrame(data.shape, columns = ["Shape"])
var_1
# -----

```

Columns Names

Columns

0	a
1	b
2	c
3	d
4	e
5	f
6	target

Information Dataset

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 7 columns):
 #   Column  Non-Null Count  Dtype  
---  -- 
 0   a        10000 non-null   float64
 1   b        10000 non-null   float64
 2   c        10000 non-null   float64
 3   d        10000 non-null   float64
 4   e        10000 non-null   float64
 5   f        10000 non-null   float64
 6   target   10000 non-null   float64
dtypes: float64(7)
memory usage: 547.0 KB
*****
```

Data Types from Dataset

```

a        float64
b        float64
c        float64
d        float64
e        float64
f        float64
target   float64
dtype: object

```

```
*****
```

Description

```
*****
```

	a	b	c	d	e	f
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	4.941596	10.090598	0.000504	5.488420	-0.062724	1.509434
std	2.876301	5.785891	2.867738	2.600951	5.785684	0.865024
min	0.000116	0.003155	-4.999519	1.000050	-9.999665	0.000025
25%	2.463289	5.078916	-2.462521	3.221715	-5.114556	0.770869
50%	4.925286	10.117936	0.020681	5.500272	-0.125656	1.518271
75%	7.400063	15.129584	2.446739	7.724046	5.009541	2.260338
max	9.997177	19.998497	4.999010	9.998104	9.999443	2.999819

```
*****
```

Missing Values

```
*****
```

```
a      0  
b      0  
c      0  
d      0  
e      0  
f      0  
target  0  
dtype: int64
```

```
*****
```

Duplicated Values

```
*****
```

```
0      False  
1      False  
2      False  
3      False  
4      False  
...  
9995   False  
9996   False  
9997   False  
9998   False  
9999   False  
Length: 10000, dtype: bool
```

```
*****
```

Shape of the Dataset

```
*****
```

```
Out[8]: Shape  
0 10000  
1    7
```

✓ EDA Completed

The **Exploratory Data Analysis (EDA)** has been completed successfully. The dataset do **not contain any null or missing values**, and **no duplicate records** were found.

This confirms that the dataset is **clean and well-prepared** for further analysis. We will now proceed to the **Data Visualization** step to better understand patterns and relationships within the data.

Start Data Visualization

In this step, we begin the **data visualization** process to explore patterns, trends, and relationships within the dataset.

```
In [9]: #=====
#                                         Line plot
# =====
features = ['a', 'b', 'c', 'd', 'e', 'f']
# Loop through each feature
for feature in features:
    plt.figure(figsize=(7,4))
    plt.plot(data[feature].head(10), marker='+', markeredgecolor='red', line
    # Title and Labels
    plt.title(f"Line Plot of Feature '{feature}'", fontsize=14, fontweight='
    plt.xlabel("Index")
    plt.ylabel("Value")
    plt.grid(True) # optional: adds grid
    plt.show()
#=====
#                                         Bar Plot from dataset
#=====

print("  ")
print("\033[1m" + "Bar Plot ".center(70) + "\033[0m")
print("  ")

#=====
for feature in features:
    plt.figure(figsize=(7,4))
    values = data[feature].head(10)
    x = range(len(values))
    plt.bar(x, values)
    # Title and Labels
    plt.title(f"Line Plot of Feature '{feature}'", fontsize=14, fontweight='
    plt.xlabel("Index")
    plt.ylabel("Feature")
    plt.grid(True) # optional: adds grid
    plt.show()
#=====
#                                         Histogram
#=====

print("  ")
print("\033[1m" + "Histogram".center(70) + "\033[0m")
print("  ")
plt.figure(figsize=(8,5))
plt.hist(data['a'].head(19), bins=10, color='skyblue', edgecolor='black') #
```

```

plt.title("Histogram of Feature 'a'", fontsize=14, fontweight='bold')
plt.xlabel("Value")
plt.ylabel("Frequency")
plt.grid(axis='y')
plt.show()
#=====

print("  ")
print("\u033[1m" + "Box Plot".center(70) + "\u033[0m")
print("  ")

#=====          Box Plot
#=====

plt.figure(figsize=(10,6))
# Boxplot
plt.boxplot([data[feature] for feature in features], labels=features, patch_
            boxprops=dict(facecolor='skyblue', color='black'),
            whiskerprops=dict(color='black'),
            capprops=dict(color='black'),
            medianprops=dict(color='red', linewidth=2),
            flierprops=dict(marker='o', markerfacecolor='red', markersize=6,
            plt.title("Box Plot of Features 'a' to 'f'", fontweight='bold')
            plt.xlabel("Features", fontsize=12)
            plt.ylabel("Value", fontsize=12)
            plt.grid(axis='y', alpha=0.75)
            plt.show()
#=====

print("  ")
print("\u033[1m" + "Heatmap ".center(70) + "\u033[0m")
print("  ")

#=====          Heatmap from dataset
#=====

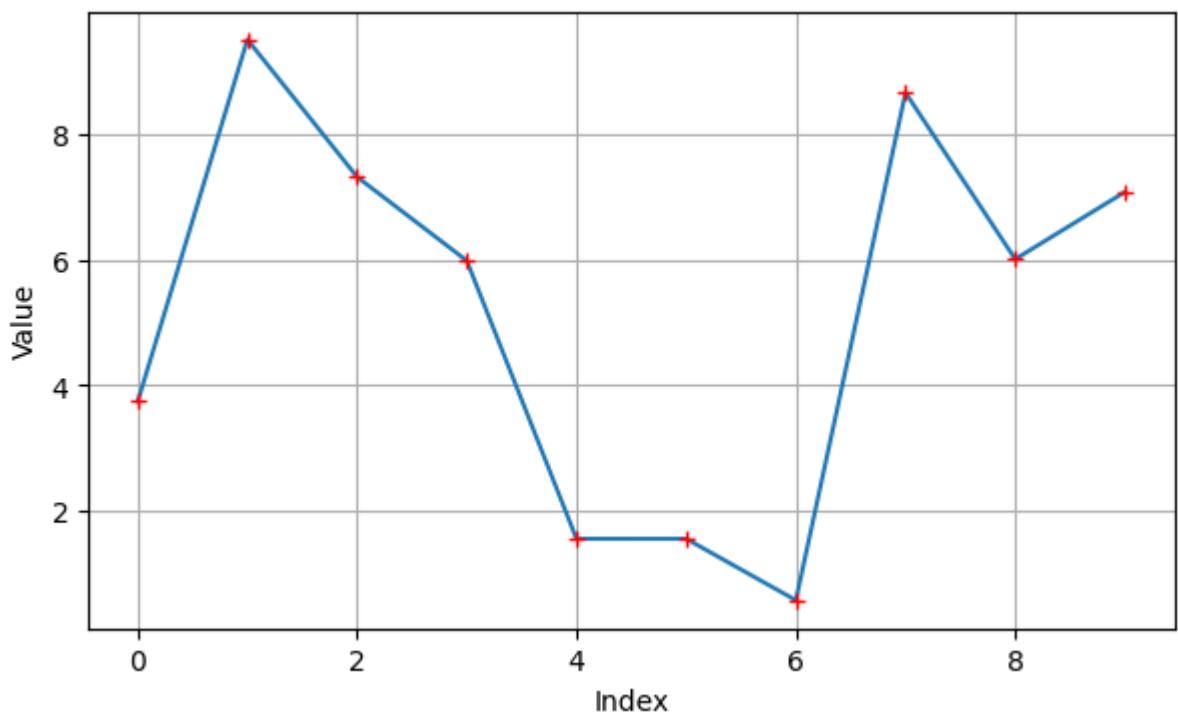
corr_matrix = data[features].corr()

# Plot heatmap
plt.figure(figsize=(8,6))
sns.heatmap(corr_matrix, annot=True, cmap="coolwarm", linewidths=1, linecolor="black",
            plt.title("Heatmap of Feature Correlations ('a' to 'f')", fontsize=14, fontweight='bold')
            plt.show()
#=====          Pair Plot /Kde plot
#=====

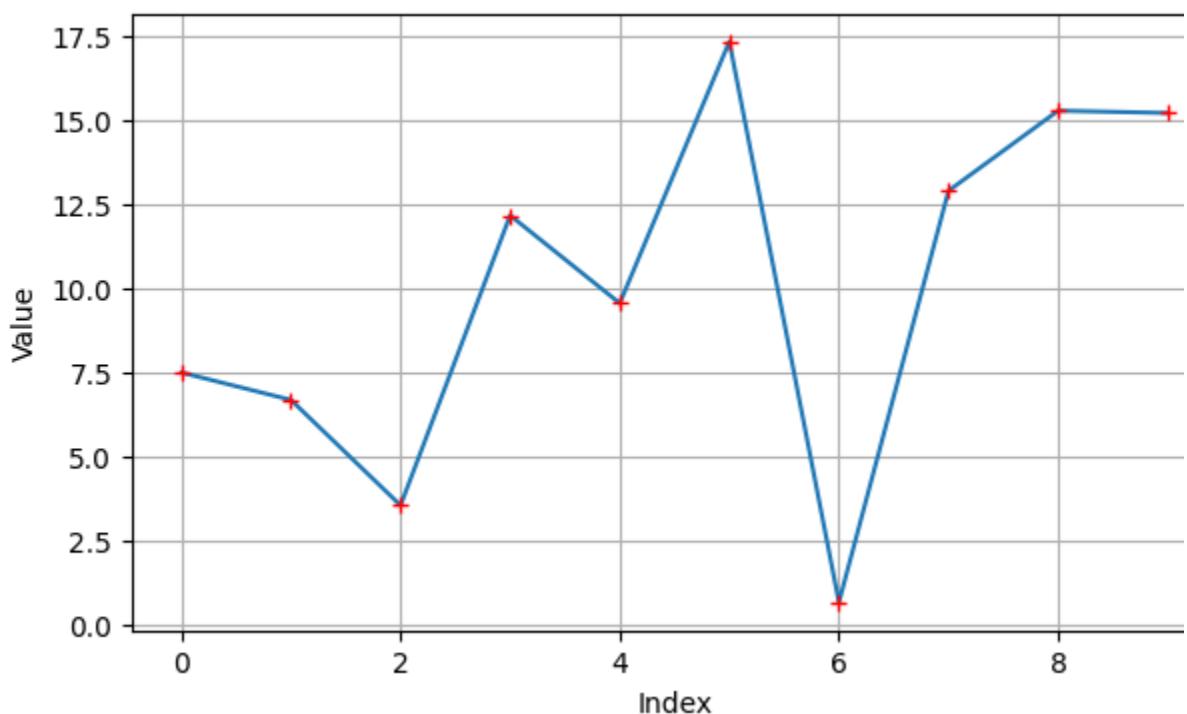
plt.figure(figsize=(8,6))
sns.kdeplot(data["a"], fill="green")
plt.show()

```

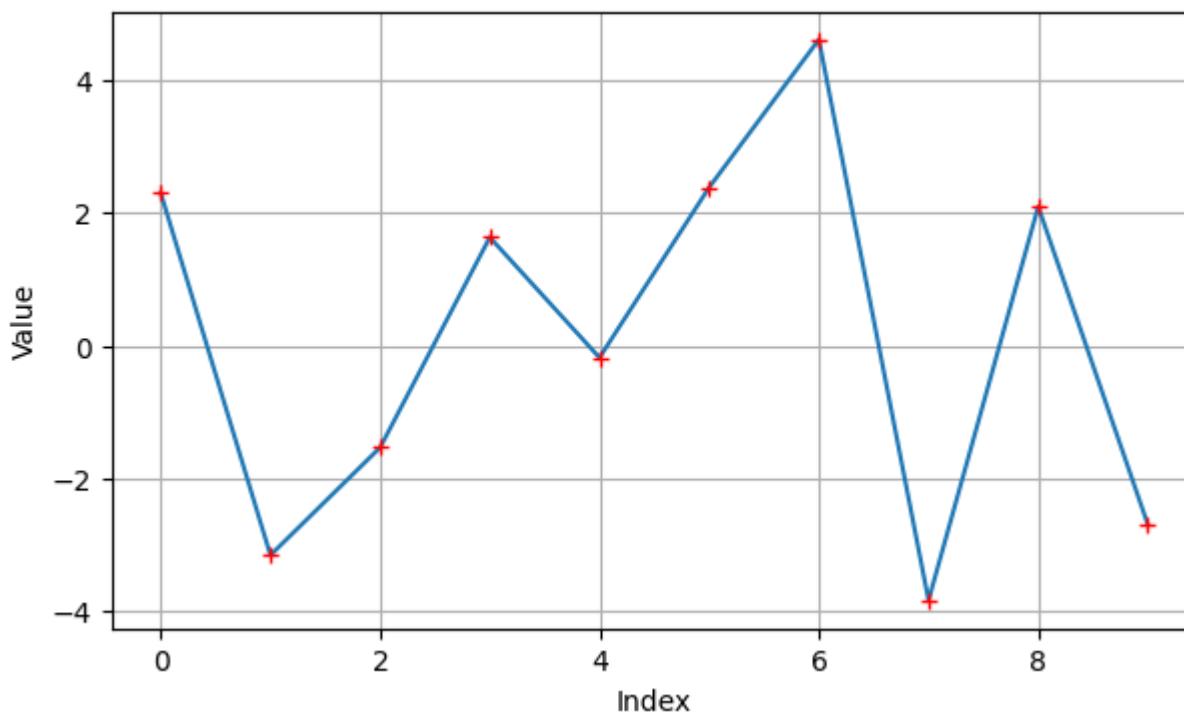
Line Plot of Feature 'a'



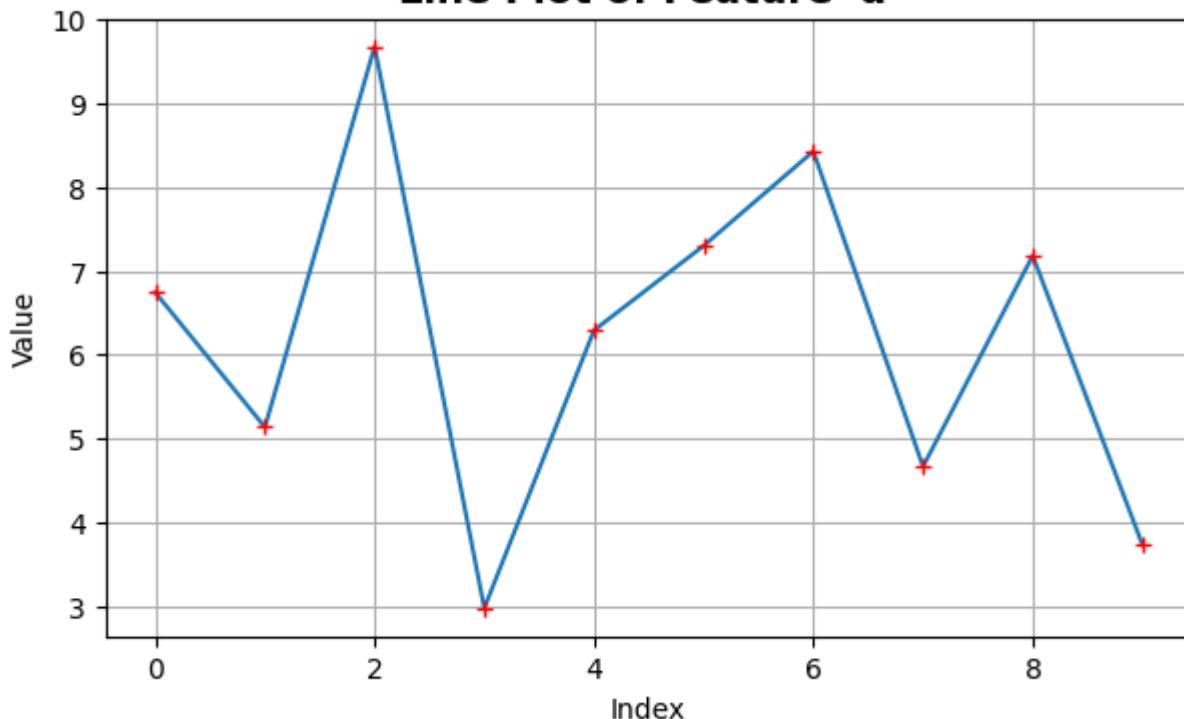
Line Plot of Feature 'b'



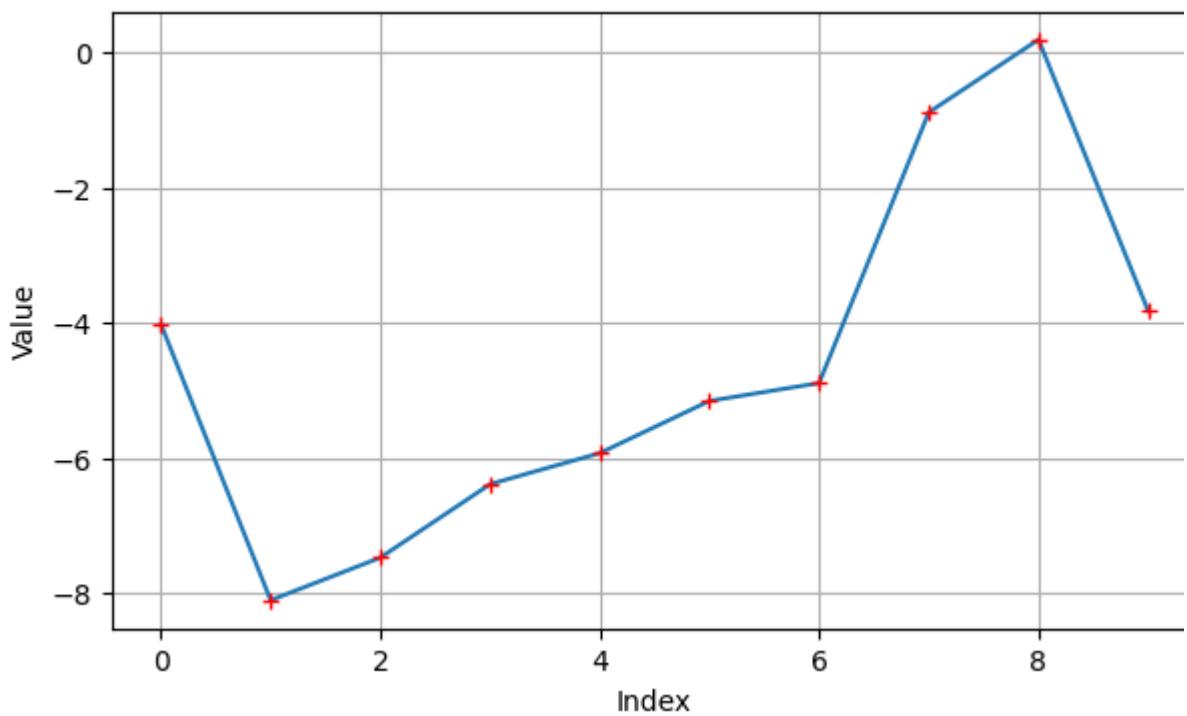
Line Plot of Feature 'c'



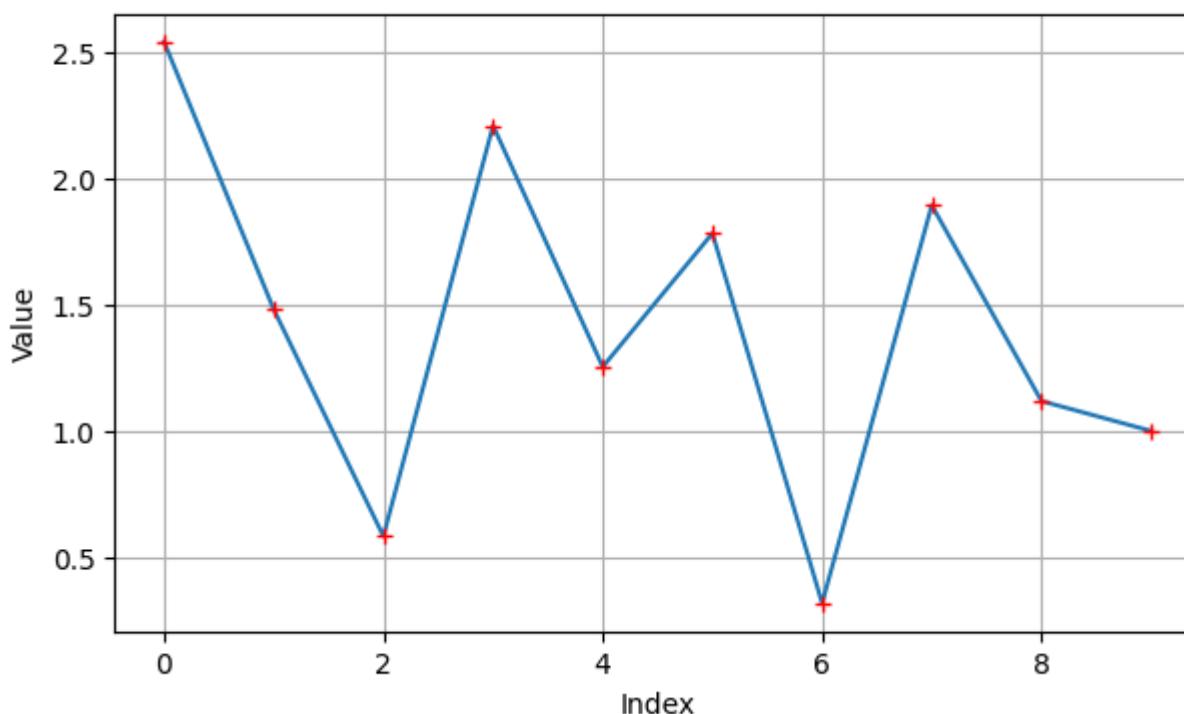
Line Plot of Feature 'd'



Line Plot of Feature 'e'

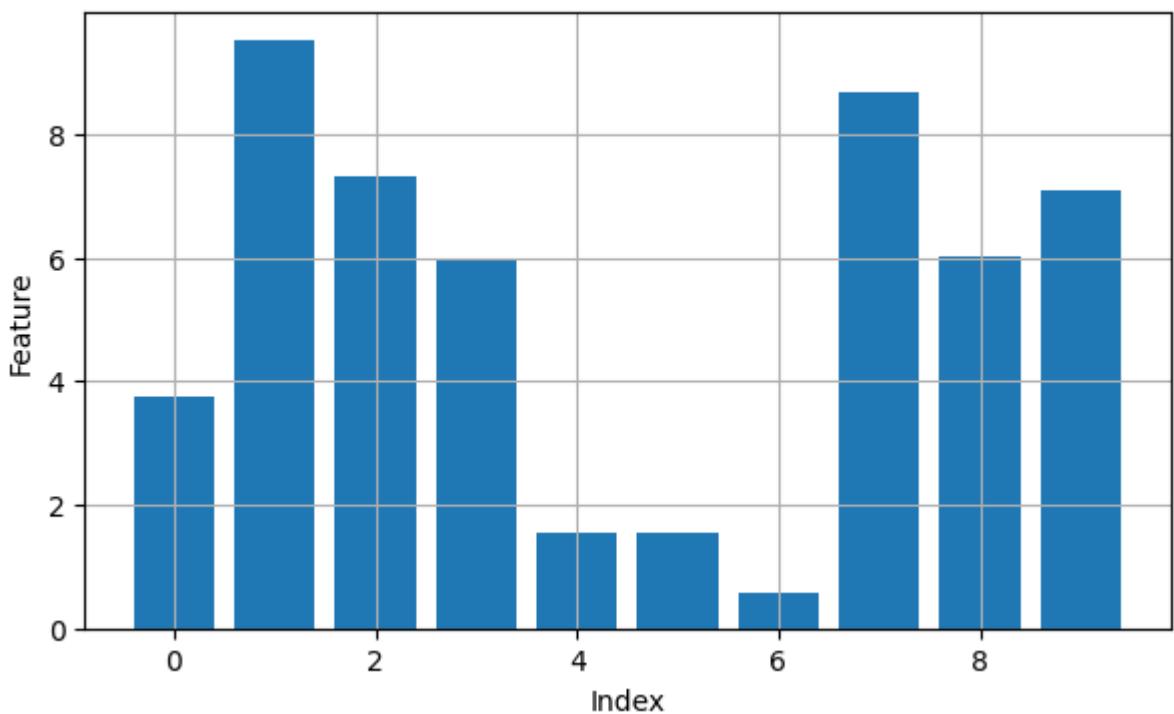


Line Plot of Feature 'f'

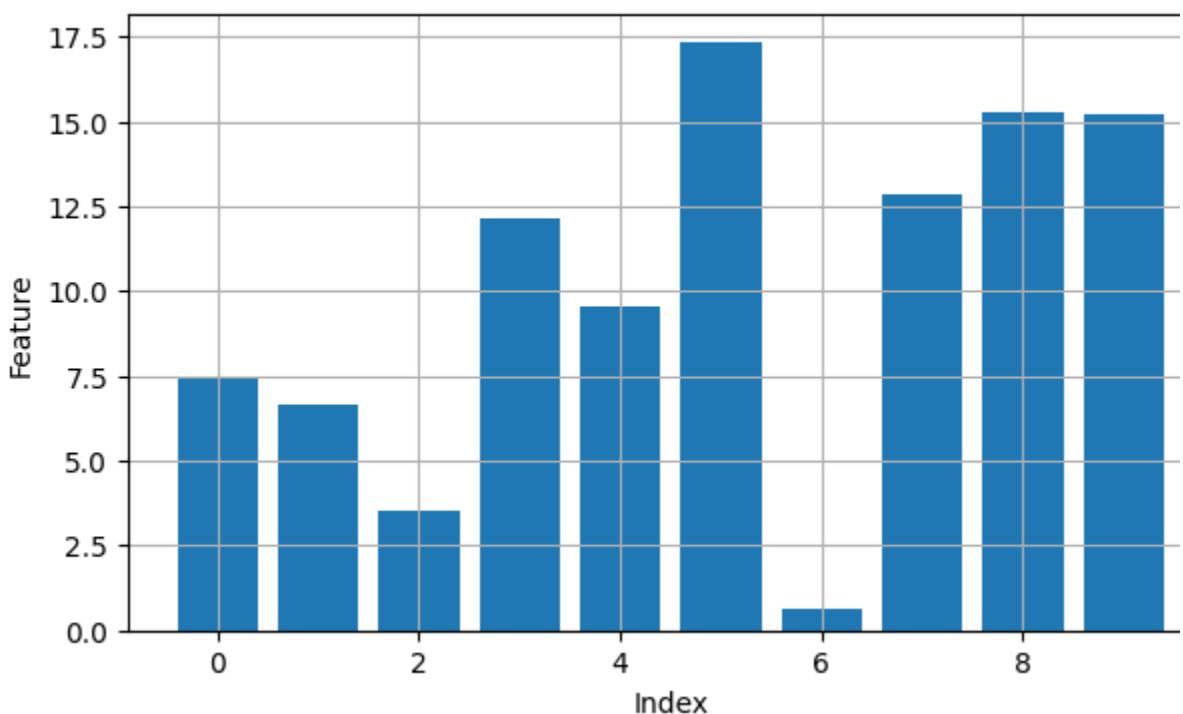


Bar Plot

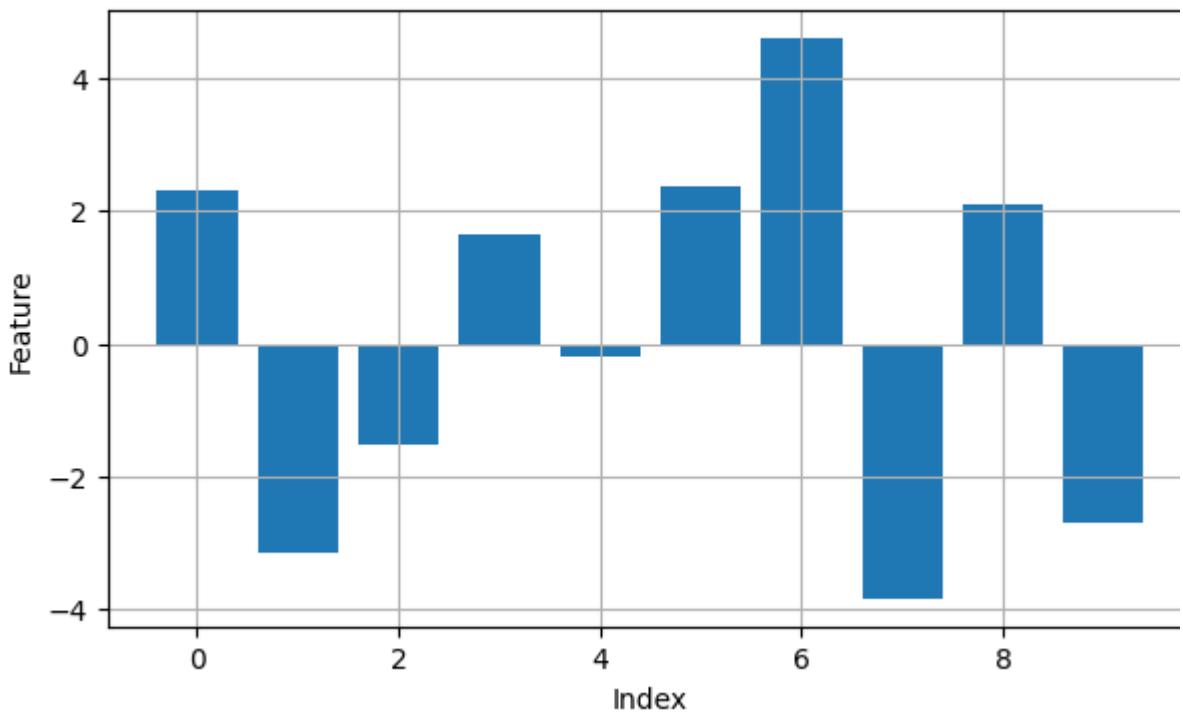
Line Plot of Feature 'a'



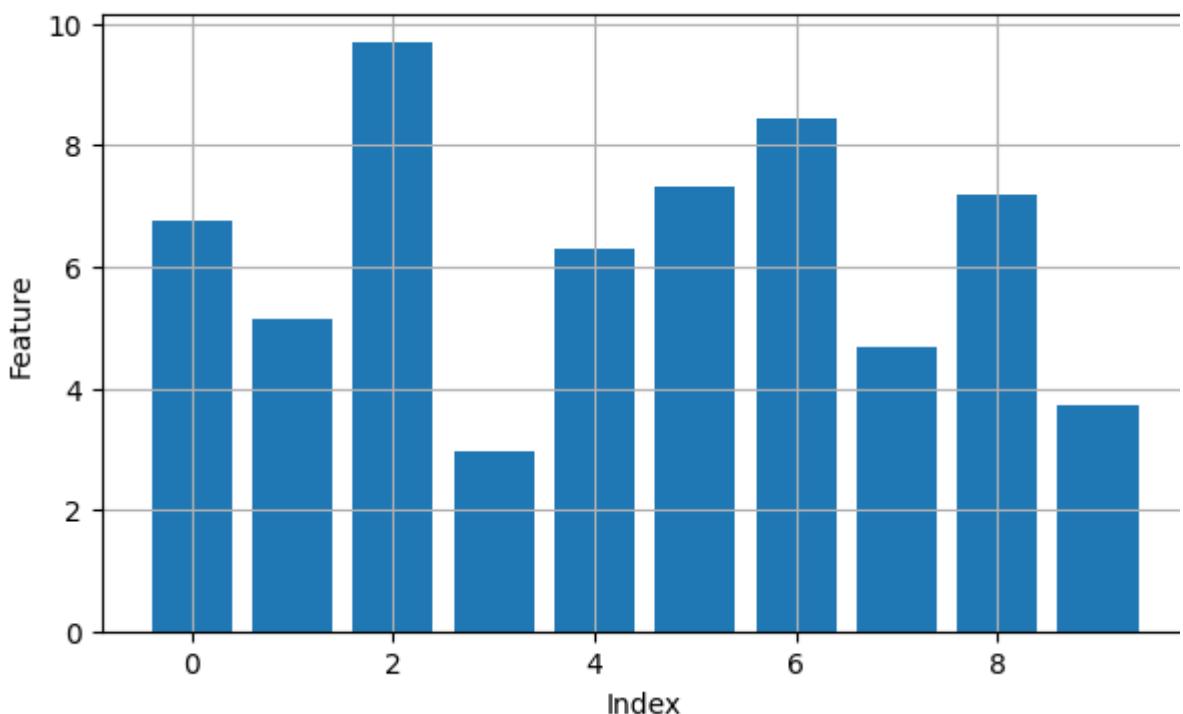
Line Plot of Feature 'b'



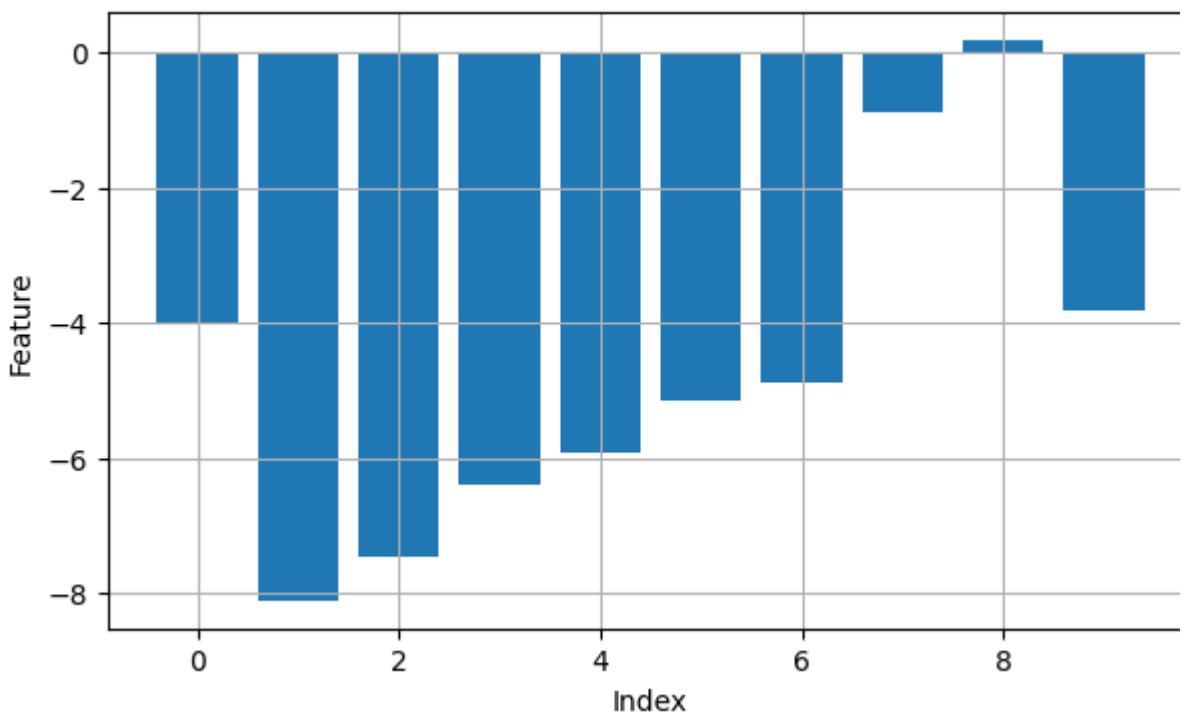
Line Plot of Feature 'c'



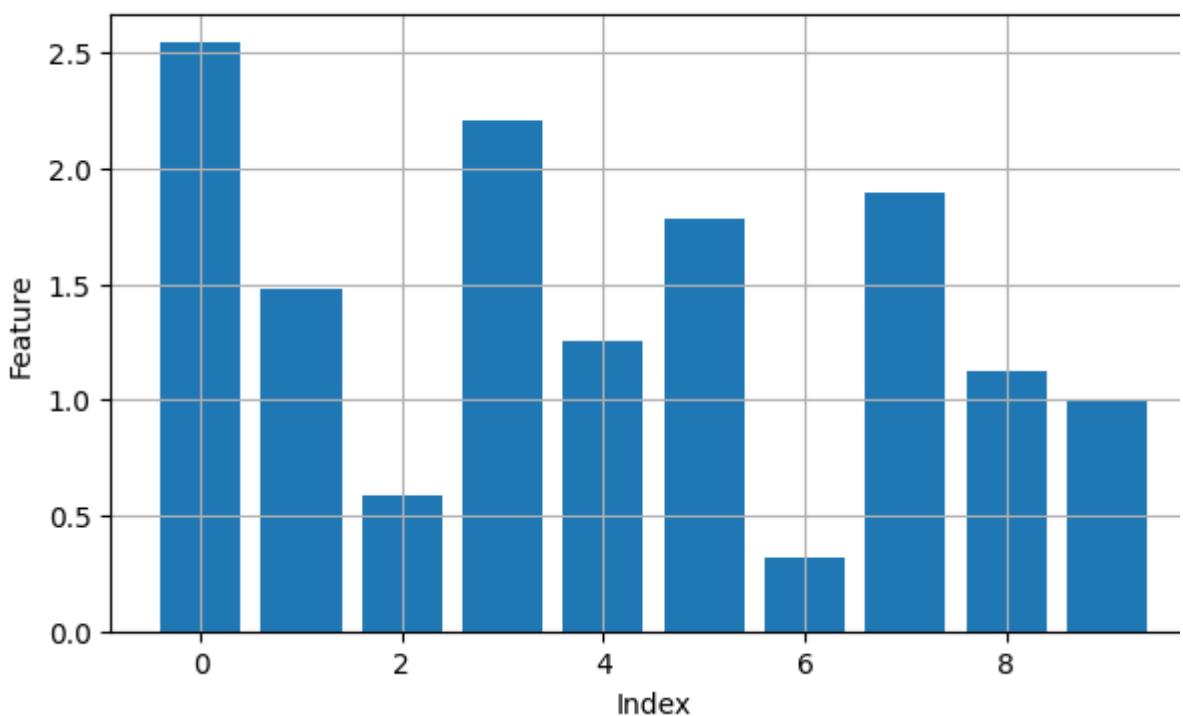
Line Plot of Feature 'd'



Line Plot of Feature 'e'

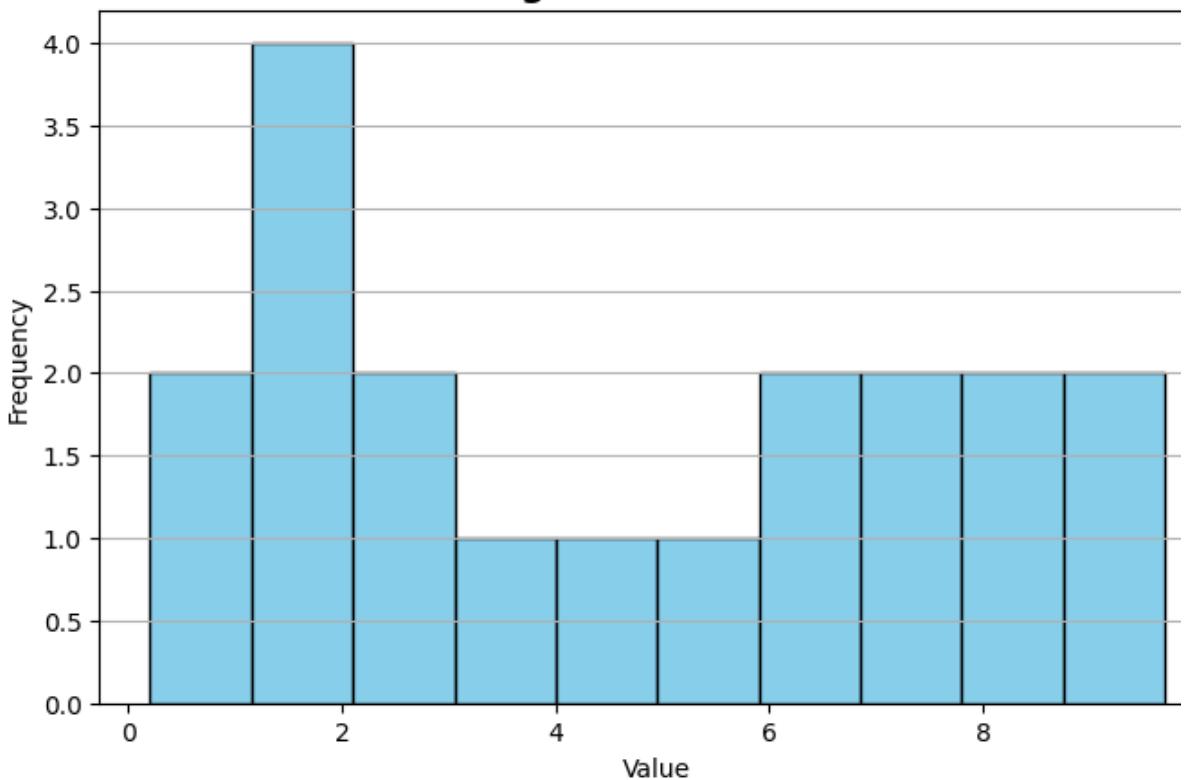


Line Plot of Feature 'f'



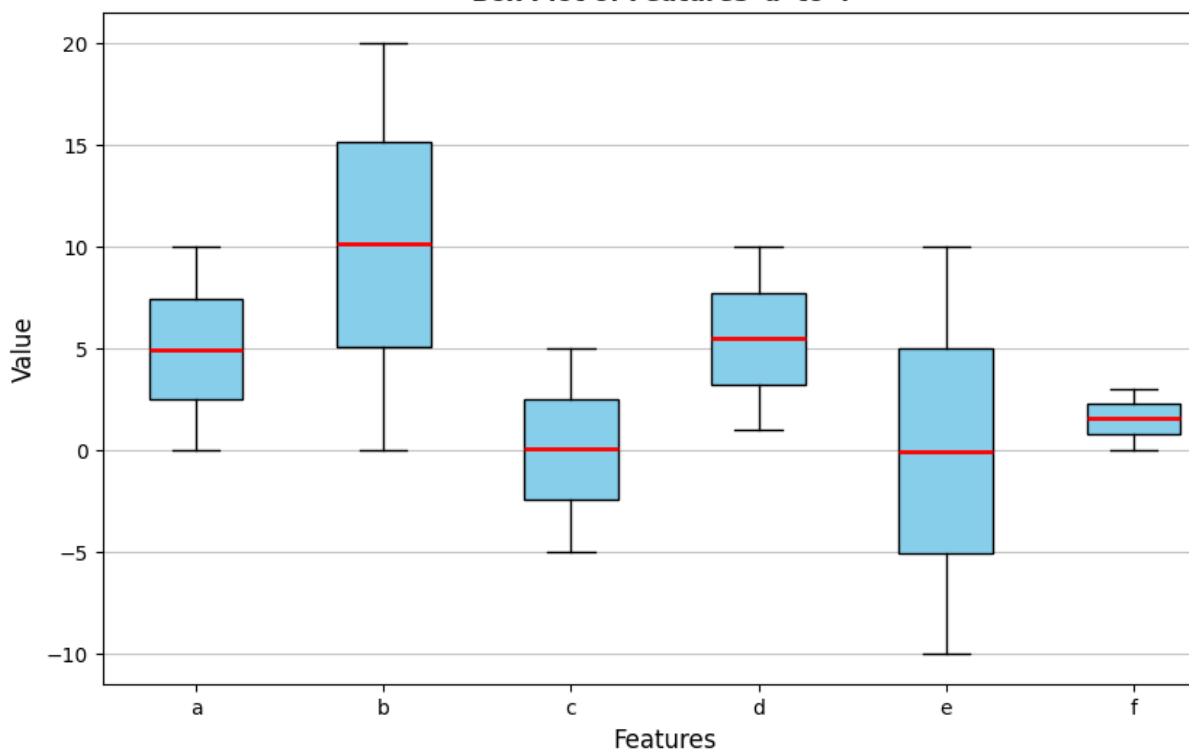
Histogram

Histogram of Feature 'a'



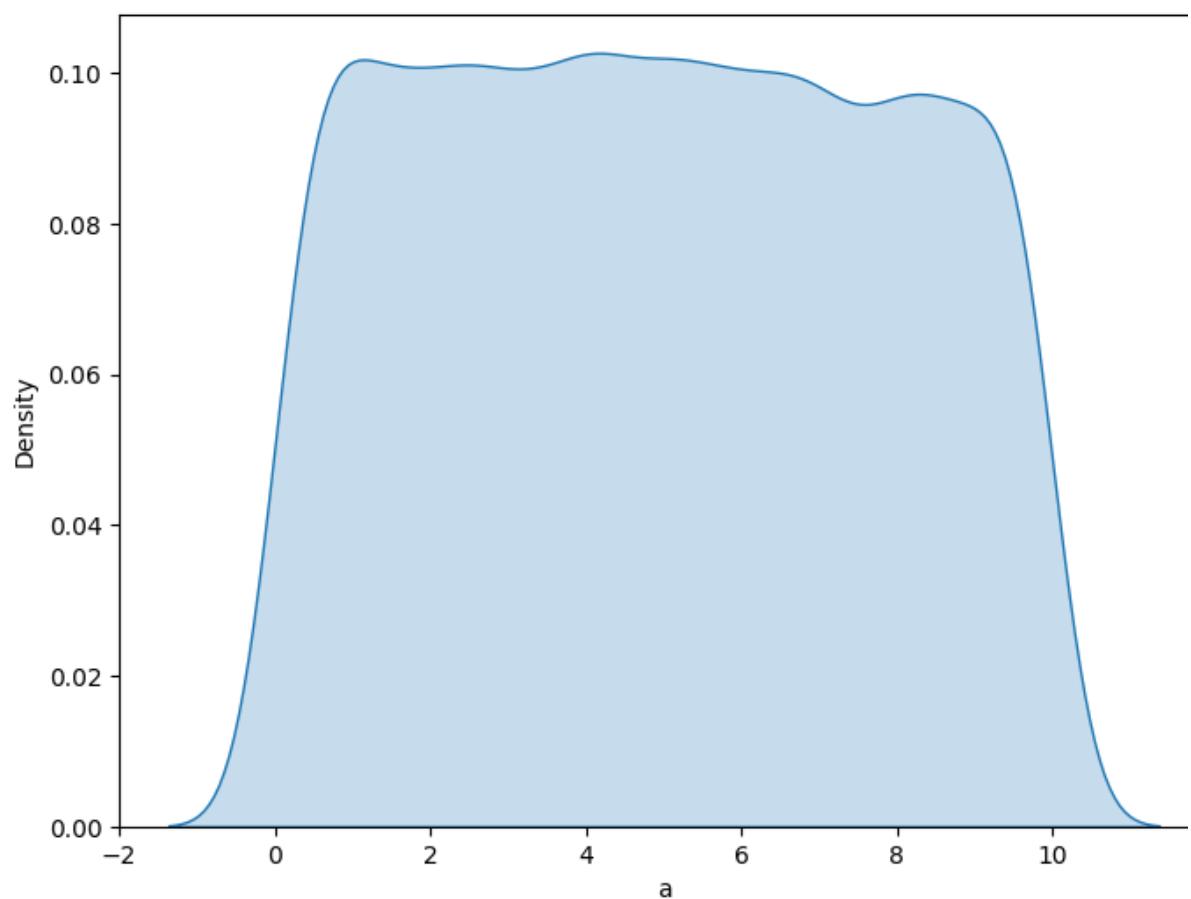
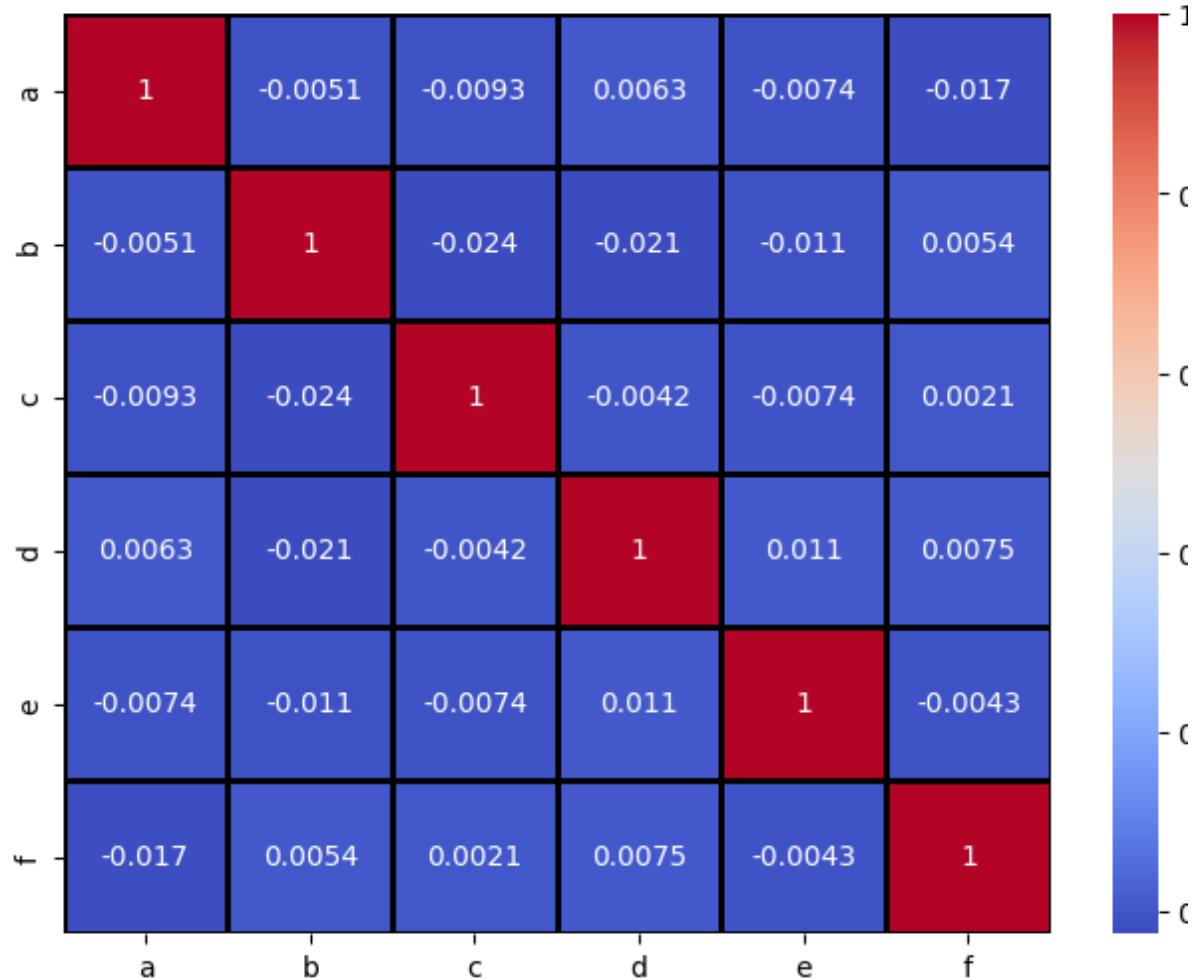
Box Plot

Box Plot of Features 'a' to 'f'



Heatmap

Heatmap of Feature Correlations ('a' to 'f')



Insight

Feature a shows a moderate spread with a few outliers. Feature d has a narrow distribution indicating consistent values. Features b and f have higher variability and multiple outliers, which need further investigation.

Feature Engineering

1. Train Test split
2. StandardScaler
3. LinearRegression
4. DecisionTreeRegressor
5. RandomForestRegressor
6. Pipeline

In [14]:

```
#=====
#                               Train Test Split
#=====

from sklearn.model_selection import train_test_split
X = data.iloc[:, :-1]
y = data.iloc[:, -1]

X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    random_state = 42,
    test_size = 0.2
)
#=====
#                               X_train_describe
#=====

print("*"*80)
print("\033[1m" + "Before Describe mean/std".center(70) + "\033[0m")
print("*"*80)
display(X_train.describe().round(2))
#=====
#                               Standard_Scalar
#=====

print("*"*80)
print("\033[1m" + "After Describe Mean/Standard derivation".center(70) + "\033[0m")
print("\033[1m" + "Used Standard Scaler".center(70) + "\033[0m")
print("*"*80)
#=====

#Import Model
#=====

from sklearn.preprocessing import StandardScaler
Scaler = StandardScaler()
#=====

#Fit/Transform model
#=====

Scaler.fit(X_train)
```

```

X_train_scaled = Scaler.transform(X_train)
X_test_scaled = Scaler.transform(X_test)
X_train_Scaler = pd.DataFrame(X_train_scaled, columns = X_train.columns)
X_test_Scaler = pd.DataFrame(X_test_scaled, columns = X_train.columns)
#=====
# Describe
#=====

X_test_Scaler.describe()
print("The StandardScaler model is working correctly because the mean is 0 and the standard deviation is 1.")


*****

```

Before Describe mean/std

	a	b	c	d	e	f
count	8000.00	8000.00	8000.00	8000.00	8000.00	8000.00
mean	4.93	10.07	0.00	5.49	-0.07	1.51
std	2.88	5.80	2.87	2.60	5.78	0.86
min	0.00	0.00	-5.00	1.00	-10.00	0.00
25%	2.44	5.03	-2.46	3.23	-5.12	0.77
50%	4.92	10.06	0.01	5.52	-0.12	1.52
75%	7.38	15.17	2.44	7.73	4.99	2.26
max	10.00	20.00	5.00	10.00	10.00	3.00

After Describe Mean/Standard derivation

Used Standard Scaler

The StandardScaler model is working correctly because the mean is 0 and the standard deviation is 1.

In [44]:

```

#=====
#                               Linear Regression
#=====

#=====
#Import model
#=====

from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_squared_error
#=====

# Fit/transfrom model
#=====

print("*"*50)
print("That is linear Regression Score R2:")
print("*"*50)
lr_model = LinearRegression()
lr_model.fit(X_train_scaled, y_train)
y_pred = lr_model.predict(X_test_scaled)
r2 = r2_score(y_test, y_pred)
rmse = mean_squared_error(y_test, y_pred)

print("R2 Score:", r2)
print("RMSE:", rmse)
#=====

#DecisionTreeRegressor
#=====

from sklearn.tree import DecisionTreeRegressor

```

```

print(" ")
print("*"*50)
print("That is DecisionTreeRegressor Score R2:")
print("*"*50)
for depth in [10]:
    model = DecisionTreeRegressor(max_depth=depth, random_state=42)
    #=====
    #Fit/Predict model
    #=====
    model.fit(X_train, y_train)
    pred = model.predict(X_test)
    print("R2:", r2_score(y_test, pred))

#=====
# RandomForestRegressor
#=====

from sklearn.ensemble import RandomForestRegressor

rf = RandomForestRegressor(
    n_estimators=200, ## 200 trees
    max_depth=10, # maximum height 10 levels
    min_samples_split=5, # model is stable
    min_samples_leaf=2, # Overfitting control
    random_state=42 # Data randomly split
)
#=====
# Fit/Predict Model
#=====

rf.fit(X_train, y_train)
y_pred = rf.predict(X_test)
print(" ")
print("*"*50)
print("That is RandomForestRegressor Score:")
print("*"*50)
print("R2:", r2_score(y_test, y_pred))
print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred)))

```

That is linear Regression Score R2:

R2 Score: 0.5281134635297018

RMSE: 2509.0934429103254

That is DecisionTreeRegressor Score R2:

R2: 0.9046040055320388

That is RandomForestRegressor Score:

R2: 0.9628245221688397

RMSE: 14.0594361629072

from IPython.display import display, HTML

display(HTML("""

Model Performance Conclusion

◆ Linear Regression

- **R² Score:** 0.53
- **RMSE:** 2509.09

Linear Regression shows **poor performance** on this dataset. This indicates that the relationship between features and target is **not purely linear**. The high RMSE value also suggests large prediction errors.

◆ Decision Tree Regressor

- **R² Score:** 0.90

Decision Tree Regressor significantly improves performance by capturing **non-linear patterns** in the data. This model provides a strong fit with much better explanatory power than Linear Regression.

◆ Random Forest Regressor ★ (Best Model)

- **R² Score:** 0.96
- **RMSE:** 14.05

Random Forest Regressor delivers the **best performance** among all models. With the highest R² score and the lowest RMSE, it demonstrates excellent generalization and accuracy. The ensemble approach reduces overfitting and captures complex feature interactions effectively.

Final Conclusion

Based on the evaluation metrics, **Random Forest Regressor** is the most suitable model for this dataset. It outperforms Linear Regression and Decision Tree by a significant margin and is recommended for **real-world deployment** and **future predictions**.



In []: